Multiple Linear Regression - Exercise

You are given a real estate dataset.

Real estate is one of those examples that every regression course goes through as it is extremely easy to understand and there is a (almost always) certain causal relationship to be found.

The data is located in the file: 'real_estate_price_size_year.csv'.

You are expected to create a multiple linear regression (similar to the one in the lecture), using the new data.

In this exercise, the dependent variable is 'price', while the independent variables are 'size' and 'year'.

Good luck!

Import the relevant libraries

In [1]:
```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn
seaborn.set()
```

Load the data

In [2]:
```
data = pd.read_csv('real_estate_price_size_year.csv')
```

In [3]:
```
data
```

Out[3]:

	price	size	year
0	234314.144	643.09	2015
1	228581.528	656.22	2009
2	281626.336	487.29	2018
3	401255.608	1504.75	2015
4	458674.256	1275.46	2009
...	...	...	...
95	252460.400	549.80	2009
96	310522.592	1037.44	2009
97	383635.568	1504.75	2006
98	225145.248	648.29	2015
99	274922.856	705.29	2006

100 rows × 3 columns

In [4]:
```
data.head()
```

Out[4]:

	price	size	year
0	234314.144	643.09	2015
1	228581.528	656.22	2009
2	281626.336	487.29	2018
3	401255.608	1504.75	2015
4	458674.256	1275.46	2009

In [5]:
```
data.describe()
```

Out[5]:

	price	size	year
count	100.000000	100.000000	100.000000
mean	292289.470160	853.024200	2012.600000
std	77051.727525	297.941951	4.729021
min	154282.128000	479.750000	2006.000000
25%	234280.148000	643.330000	2009.000000
50%	280590.716000	696.405000	2015.000000
75%	335723.696000	1029.322500	2018.000000
max	500681.128000	1842.510000	2018.000000

Create the regression

Declare the dependent and the independent variables

In [7]:
```
y = data['price']
x1 = data[['size', 'year']]
```

Regression

In [8]:
```
x = sm.add_constant(x1)
results = sm.OLS(y,x).fit()
```

In [9]:
```
results.summary()
```

Out[9]:

OLS Regression Results

Dep. Variable:	price	**R-squared:**	0.776
Model:	OLS	**Adj. R-squared:**	0.772
Method:	Least Squares	**F-statistic:**	168.5
Date:	Wed, 25 Aug 2021	**Prob (F-statistic):**	2.77e-32
Time:	22:06:01	**Log-Likelihood:**	-1191.7
No. Observations:	100	**AIC:**	2389.
Df Residuals:	97	**BIC:**	2397.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
const	-5.772e+06	1.58e+06	-3.647	0.000	-8.91e+06	-2.63e+06
size	227.7009	12.474	18.254	0.000	202.943	252.458
year	2916.7853	785.896	3.711	0.000	1357.000	4476.571

Omnibus:	10.083	**Durbin-Watson:**	2.250
Prob(Omnibus):	0.006	**Jarque-Bera (JB):**	3.678
Skew:	0.095	**Prob(JB):**	0.159
Kurtosis:	2.080	**Cond. No.**	9.41e+05

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 9.41e+05. This might indicate that there are
strong multicollinearity or other numerical problems.

In [ ]:
```

```